

In [108]:

```
import pandas as pd
import numpy as np
from sklearn import decomposition
import matplotlib.pyplot as plt
```

In [109]:

```
data = pd.read_csv("rool_data.csv")
```

In [110]:

```
y= data.iloc[:, -1].values
```

In [74]:

```
data.head()
```

Out[74]:

	Acedamic percentage in Operating Systems	percentage in Algorithms	Percentage in Programming Concepts	Percentage in Software Engineering	Percentage in Computer Networks	Percentage in Electronics Subjects	Percentage in Computer Architecture	I
0	69	63	78	87	94	94	87	
1	78	62	73	60	71	70	73	
2	71	86	91	87	61	81	72	
3	76	87	60	84	89	73	62	
4	92	62	90	67	71	89	73	

5 rows × 39 columns

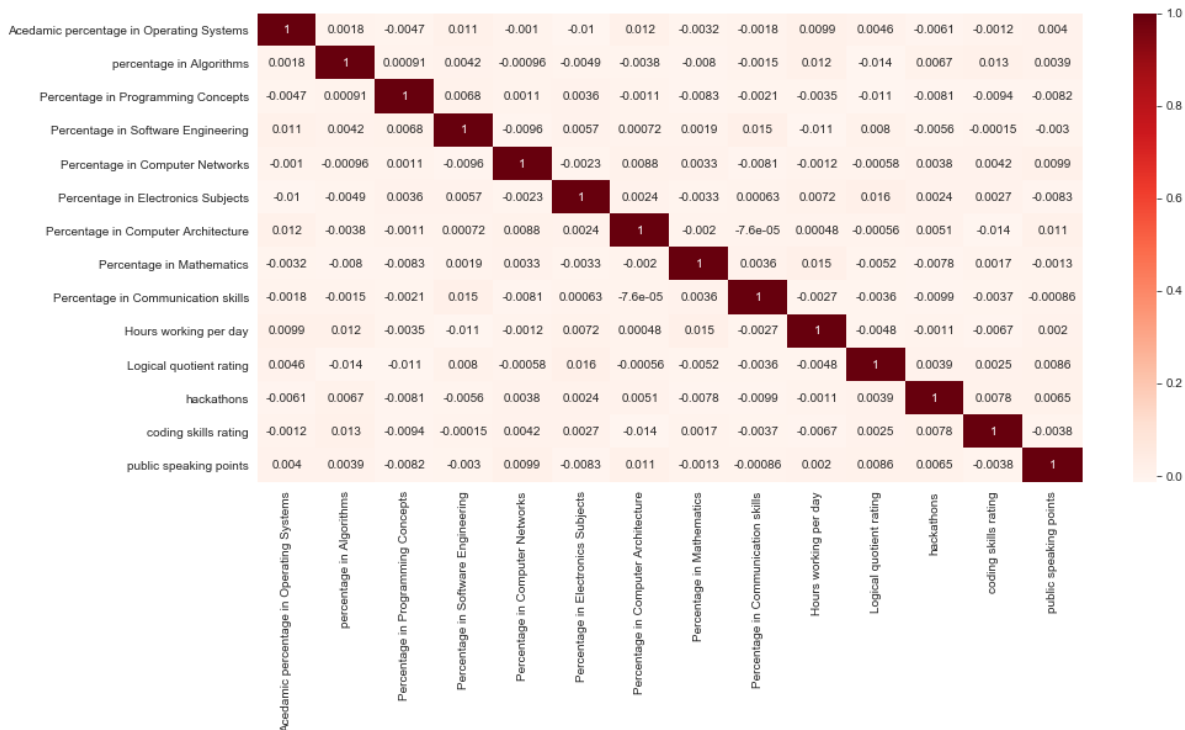
In [75]:

```
import seaborn as sns
import matplotlib.pyplot as plt
sns.set_style('whitegrid')
pd.options.display.float_format='{:.3f}'.format
```

Co-relation heat Map

In [76]:

```
corrMatrix=data.corr()
plt.figure(figsize=(15,7))
sns.heatmap(corrMatrix,annot=True,cmap='Reds')
plt.show()
```



PCA

In [77]:

```
from sklearn.decomposition import PCA
```

In [78]:

```
pca_set = PCA(n_components=2)
```

In [88]:

```
Xdata = data.iloc[:, :-1].values
```

In [89]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

In [90]:

```
labelencoder = LabelEncoder()
```

In [91]:

```
for i in range(14,38):
    Xdata[:,i] = labelencoder.fit_transform(Xdata[:,i])
```

In [92]:

```
component_set = pca_set.fit_transform(Xdata)
```

In [93]:

```
component_set
```

Out[93]:

```
array([[ 6.92455735, 15.51628793],
       [-9.64836184, -1.08584709],
       [22.83910281, -12.37537294],
       ...,
       [11.28384035, -15.06111604],
       [ 6.34199002,  2.98779088],
       [-1.00708864, -15.33193353]])
```

In []:

In [94]:

```
X1=data[['Acedamic percentage in Operating Systems','percentage in Algorithms','Perceptron']]
```

In [95]:

```
component_set = pca_set.fit_transform(X1)
```

In [96]:

```
principal_set = pd.DataFrame(data = component_set
                             , columns = ['p1', 'p2'])
```

In [97]:

```
principal_set.head()
```

Out[97]:

	p1	p2
0	6.465	15.272
1	-9.834	-1.440
2	23.364	-11.777
3	-13.178	-2.327
4	8.394	8.728

In [98]:

```
final_pca = pd.concat([principal_set, data[['Suggested Job Role']]], axis = 1)
```

In [99]:

```
final_pca.head()
```

Out[99]:

	p1	p2	Suggested Job Role
0	6.465	15.272	Database Developer
1	-9.834	-1.440	Portal Administrator
2	23.364	-11.777	Portal Administrator
3	-13.178	-2.327	Systems Security Administrator
4	8.394	8.728	Business Systems Analyst

PCA visualization 2D projection

In [100]:


```
fig = plt.figure(figsize = (8,8))
```

<Figure size 576x576 with 0 Axes>

In [101]:

```
fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('p1', fontsize = 15)
ax.set_ylabel('p2', fontsize = 15)
ax.set_title('2 component PCA', fontsize = 20)

targets = ['Applications Developer', 'Business Intelligence Analyst', 'Business Systems Analyst']
colors = ['r', 'g', 'b', 'rb', 'br']
for y, color in zip(targets, colors):
    indicesToKeep = final_pca['Suggested Job Role'] == y
    ax.scatter(final_pca.loc[indicesToKeep, 'p1'],
               final_pca.loc[indicesToKeep, 'p2'],
               c = color,
               s = 5)
ax.legend(targets)
ax.grid()
```



In []:

In []:

In []:

In []:

Dummies Variable

In [59]:

```
dummiesC = pd.get_dummies(data[ 'certifications' ])
```

In [60]:

```
dummiesC.head()
```

Out[60]:

	app development	distro making	full stack	hadoop	information security	machine learning	python	r programming	s programr
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	1	0	0	
2	1	0	0	0	0	0	0	0	
3	0	0	0	0	0	0	1	0	
4	1	0	0	0	0	0	0	0	

In [61]:

```
X = pd.DataFrame(data)
```

In [62]:

```
X1=data[ [ 'Acedamic percentage in Operating Systems', 'percentage in Algorithms', 'Perce
```

In [63]:

```
X1.head()
```

Out[63]:

	Acedamic percentage in Operating Systems	percentage in Algorithms	Percentage in Programming Concepts	Percentage in Software Engineering	Percentage in Computer Networks	Percentage in Electronics Subjects	Percentage in Computer Architecture	I M:
0	69	63	78	87	94	94	87	
1	78	62	73	60	71	70	73	
2	71	86	91	87	61	81	72	
3	76	87	60	84	89	73	62	
4	92	62	90	67	71	89	73	

In [64]:

```
df_dummies= pd.concat([X1,dummiesC],axis='columns')
```

In [65]:

```
df_dummies.head()
```

Out[65]:

	Acedamic percentage in Operating Systems	percentage in Algorithms	Percentage in Programming Concepts	Percentage in Software Engineering	Percentage in Computer Networks	Percentage in Electronics Subjects	Percentage in Computer Architecture	I M:
0	69	63	78	87	94	94	87	
1	78	62	73	60	71	70	73	
2	71	86	91	87	61	81	72	
3	76	87	60	84	89	73	62	
4	92	62	90	67	71	89	73	

5 rows × 23 columns

In [111]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

In [112]:

```
labelencoder_X = LabelEncoder()
```

In [116]:

```
y1 = labelencoder_X.fit_transform(y)
```

In [117]:

```
y1 = pd.DataFrame(y1)
```

In []:

Model Fitting

In [118]:

```
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
```

In [119]:

```
X_train,X_test,y_train,y_test=train_test_split(df_dummies,y1,test_size=0.2,random_st
```

Decision tree classifier

In [120]:

```
clf = tree.DecisionTreeClassifier()
clf = clf.fit(X_train, y_train)
```

In [121]:

```
from sklearn.metrics import confusion_matrix,accuracy_score
```

In [122]:

```
y_pred = clf.predict(X_test)
```

In [123]:

```
y_pred
```

Out[123]:

```
array([11,  2,  8, ..., 29, 15, 21])
```

In [124]:

```
DA=clf.score(df_dummies,y1)
DA
```

Out[124]:

```
0.8049
```

XGBClassifier

In [125]:

```
from xgboost import XGBClassifier
```

In [126]:

```
model = XGBClassifier()  
model.fit(X_train, y_train)
```

C:\Users\Ashok\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:219: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\Ashok\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:252: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Out[126]:

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
              colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=  
0,  
              max_depth=3, min_child_weight=1, missing=None, n_estimators=10  
0,  
              n_jobs=1, nthread=None, objective='multi:softprob', random_stat  
e=0,  
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
              silent=True, subsample=1)
```

In [127]:

```
y_pred = model.predict(X_test)
```

In [128]:

```
DXgbA=model.score(df_dummies,y1)  
DXgbA
```

Out[128]:

0.1495

SVM

In [129]:

```
from sklearn import svm
```

In [130]:

```
clf = svm.SVC()
clf.fit(X_train, y_train)
```

C:\Users\Ashok\Anaconda3\lib\site-packages\sklearn\utils\validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\Ashok\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

```
"avoid this warning.", FutureWarning)
```

Out[130]:

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
    kernel='rbf', max_iter=-1, probability=False, random_state=None,
    shrinking=True, tol=0.001, verbose=False)
```

In [131]:

```
y_pred = clf.predict(X_test)
```

In [132]:

```
SA=clf.score(df_dummies,y1)
SA
```

Out[132]:

```
0.8112
```

RANDOM forestRegreesor

In [133]:

```
from sklearn.ensemble import RandomForestRegressor

regressor = RandomForestRegressor(n_estimators=20, random_state=0)
regressor.fit(X_train, y_train)
y_pred = regressor.predict(X_test)
```

C:\Users\Ashok\Anaconda3\lib\site-packages\ipykernel_launcher.py:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
after removing the cwd from sys.path.
```

In [134]:

```
RFRA=regressor.score(df_dummies,y1)
RFRA
```

Out[134]:

0.6528025688186774

RANDOM FOREST CLASSIFIER

In [135]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [136]:

```
from sklearn.datasets import make_classification
```

In [137]:

```
clf = RandomForestClassifier()
```

In [138]:

```
clf.fit(X_train, y_train)
```

```
C:\Users\Ashok\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:
246: FutureWarning: The default value of n_estimators will change from
10 in version 0.20 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
C:\Users\Ashok\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: Da
taConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example us
ing ravel().
```

```
"""Entry point for launching an IPython kernel.
```

Out[138]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='g
ini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=Non
e,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

In [139]:

```
y_pred = clf.predict(X_test)
```

In [140]:

```
RFACA=clf.score(df_dummies,y1)
```

In [141]:

```
X1, y1 = make_classification(n_samples=1000, n_features=10,  
                             n_informative=2, n_redundant=0,  
                             random_state=0, shuffle=False)
```

In [142]:

```
clf = RandomForestClassifier(n_estimators=100, max_depth=2,  
                             random_state=0)  
clf.fit(X1, y1)
```

Out[142]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='g  
ini',  
                        max_depth=2, max_features='auto', max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=Non  
e,  
                        oob_score=False, random_state=0, verbose=0, warm_start=Fal  
se)
```

In [143]:

```
print(clf.feature_importances_)
```

```
[0.10709928 0.5484963 0.00847003 0.06325648 0.03568049 0.04204675  
 0.03366965 0.0590075 0.05193304 0.05034048]
```

In []:

In []:

In [144]:

```
X1[0]
```

Out[144]:

```
array([-1.66853167, -1.29901346, 0.2746472, -0.60362044, 0.7088595  
8,  
       0.42281857, -3.11685659, 0.64445203, -1.91374267, 0.6635615  
8])
```

In [145]:

```
y1[0]
```

Out[145]:

```
0
```

In [146]:

```
clf.predict([[ -1.66853167, -1.29901346,  0.2746472 , -0.60362044,  0.70885958,
              0.42281857, -3.11685659,  0.64445203, -1.91374267,  0.66356158]])
```

Out[146]:

```
array([0])
```

In [147]:

```
X2=pd.DataFrame(X1)
```

In [148]:

```
X2.head()
```

Out[148]:

	0	1	2	3	4	5	6	7	8	9
0	-1.669	-1.299	0.275	-0.604	0.709	0.423	-3.117	0.644	-1.914	0.664
1	-2.973	-1.089	-0.154	1.194	-0.098	-0.887	-0.147	1.060	0.026	-0.114
2	-0.596	-1.370	0.744	0.210	-0.006	1.366	1.555	0.613	-0.286	1.497
3	-1.069	-1.175	1.183	0.719	-1.216	0.141	-0.744	-0.159	0.240	0.100
4	-1.305	-0.966	-0.475	1.273	-1.696	0.730	-1.857	0.383	-0.887	0.878

Final Accuracy

```
Y_pred=clf.predict(X2)
```

In [150]:

```
Y_pred=clf.predict(X2)
```

In [151]:

```
finalA=clf.score(X2,Y_pred)
```

In [152]:

```
finalA
```

Out[152]:

```
1.0
```

Accuracy Comparision

all acuracies

In [155]:

```
print(DA)
print(SA)
print(RFRA)
print(RFACA)
print(DXgbA)
print(finalA)
DecisionTree_A=DA
SVM_A=SA
Regressor_A=RFRA
RandomForestClassifier_A=RFACA
XgbClassifier=DXgbA
```

```
0.8049
0.8112
0.6528025688186774
0.8032
0.1495
1.0
```

In []: